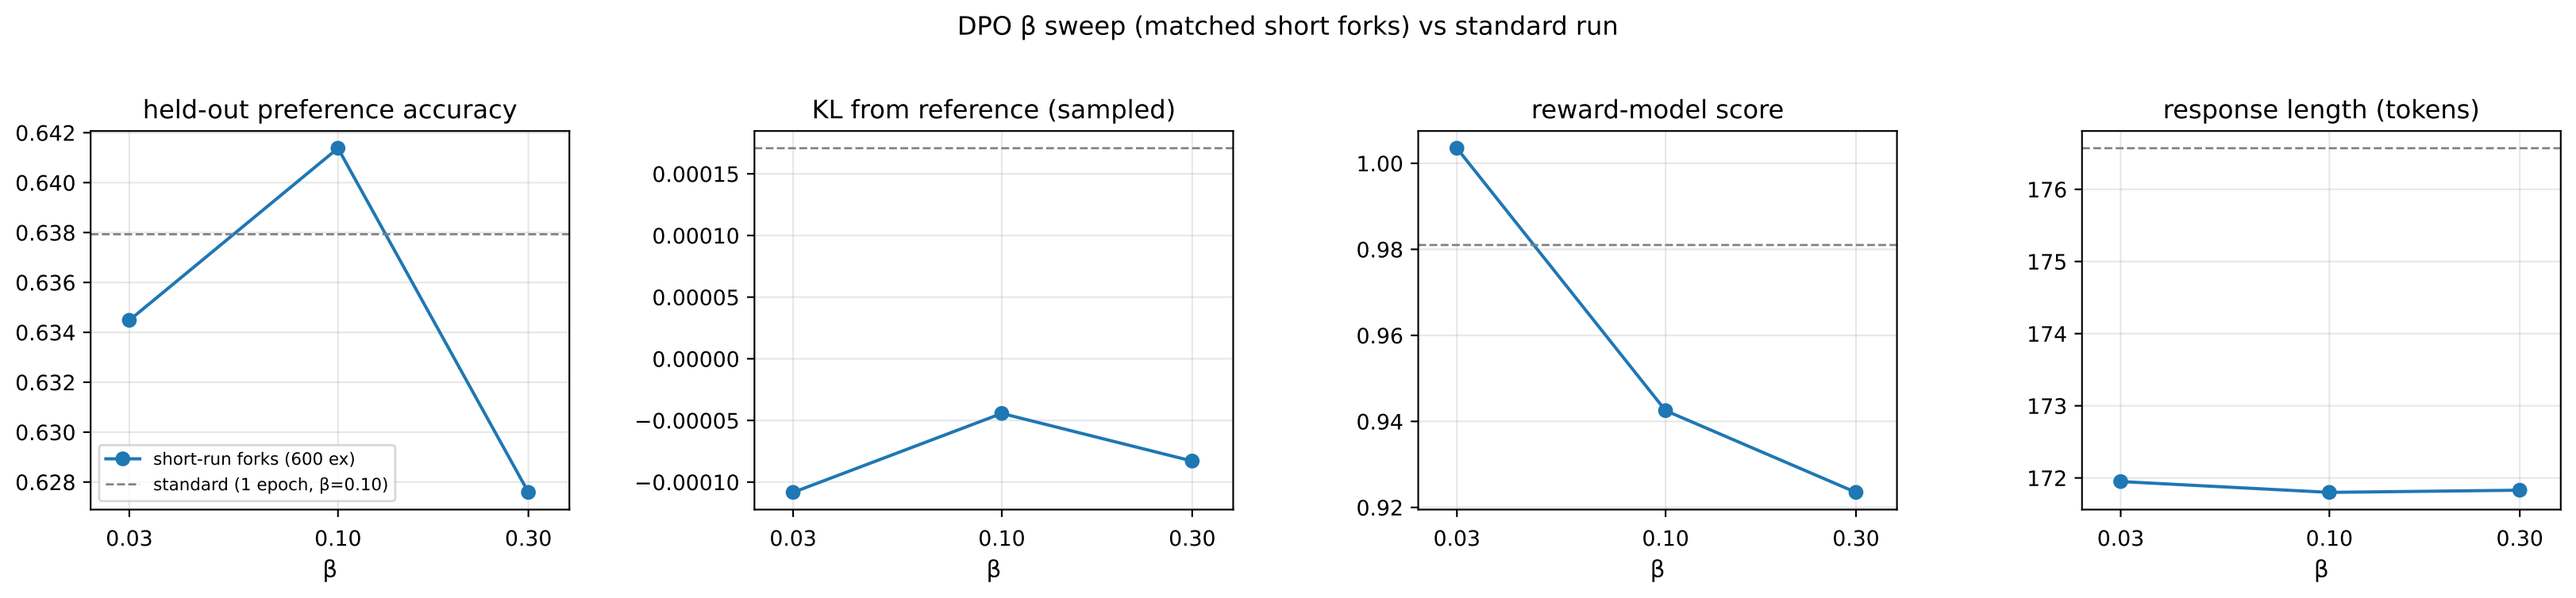
DPO β sweep (matched short forks) vs standard run
held-out preference accuracy
KL from reference (sampled)
reward-model score
response length (tokens)
short-run forks (600 ex)
standard (1 epoch, β=0.10)
0.642
0.640
0.638
0.636
0.634
0.632
0.630
0.628
0.00015
0.00010
0.00005
0.00000
−0.00005
−0.00010
1.00
0.98
0.96
0.94
0.92
176
175
174
173
172
0.03
0.10
0.30
β