



Standard PPO continuation from the supplied midpoint (20 updates)


mean learned reward




KL from reference




policy loss




value loss




entropy (sampled)



clip fraction



policy grad norm



response length (tokens)

